



Document Number	EDCS- 1150953
Based on Template	<u>EDCS-189229 Rev 20</u>
Created By	Amrik Bains
Contributors	Arvind Kansal

Universal SXGMII Interface for a Single MultiGigabit Copper Network Port

The Universal Serial Media Independent Interface for carrying single network port over a single SERDES (USXGMII) is specified in this document to meet the following requirements:

- Convey Single network ports over an USXGMII MAC-PHY interface
- Utilize a 64/66 PCS to minimize power and serial bandwidth
- Use modified 802.3by section 108.5.2.4, to add Alignment Markers to support multiple ports over single SERDES
- System Interface operates in full duplex mode only
- Ability to send PTP time stamp from PHY to MAC to improve accuracy/jitter on encrypted/non-encrypted PTP packet with MACSec is in the ASIC
- Hardware assisted auto-negotiation for all supported speeds
- Flexibility to add new features using Extension Field in pre-amble

USXGMII isn't a single protocol, but rather an architecture that allows for the definition of specific interfaces in a way that maximizes reuse and reduces risk. There are other implementation targeted at different applications and contact Cisco for details. A PHY can implement one or more options specified above based on a particular application, cost and power optimization.

Cisco Systems Intellectual Property

Cisco encourages others to adopt this specification. To the extent that Cisco has proprietary rights to information contained in this specification, any company wishing to use this specification may do so under reasonable, non-discriminatory terms, with reciprocity, to implement and fully comply with the specification.

The reasonable non-discriminatory terms are:

Cisco will not assert any claims of any patents that are owned or controlled by Cisco essential for implementing this specification against any party for making, using, selling, importing or offering for sale a product that implements the specification, provided, however, that Cisco retains the right to assert its patents (including the right to claim past royalties) against any party that asserts a patent it owns or controls (either directly or indirectly) against any product that implements this specification. Cisco retains the right to assert its patents against any product or portion thereof that is not necessary for compliance with the specification.

Royalty-bearing licenses are available to anyone who prefers that option.

For information contact:

Standards-ipr@cisco.com

Modification History

Revision	Date	Originator	Comments
2.0	11/12/2015	Amrik Bains	First official Cisco release derived from rev2.0
2.1	04/11/2016	Amrik Bains	Typo fix on page 23 for data bit on the wire. Hex value is the same as before.
2.2	05/11/2016	Amrik Bains	Remove USXGMII flow control typo's
2.3	11/01/2016	Amrik Bains	Extension Field Type typo correction to Extension Field contains Tag/Signature associated with for time-stamp to be taken by the PHY
2.4	07/30/2019	Amrik Bains	Clarification on IPG and Pre-amble processing in the PHY
2.5	08/17/23	Dale Mohlenhoff	Clarification on Intellectual Property and remove Cisco Confidential footer

Table of Contents

1	Overview	6
1.1	In-band Control and Status Signaling	9
1.1.1	Auto-neg Mechanism	9
1.1.2	Auto-negotiation Message	10
1.1.3	Packet Control Header	13
2	Implementation Specification.....	15
2.1	XGMII Mapping	15
2.2	GMII Mapping	15
2.3	MII Mapping.....	15
2.4	Pause Frame Support	15
2.5	Auto-neg Mechanism.....	15
2.5.1	Transmitting Configuration Words	16
2.6	Rate Adaptation - Replicating Transmit Bytes	16
2.6.1	Tx 10G Mode	17
2.6.2	10G Mode Sampling Received Bytes	18
2.6.3	Tx 5G Mode	18
2.6.4	Rx 5G Sampling Received Bytes	19
2.6.5	Receiving Configuration Words	19
2.6.6	Clocking	19
2.6.7	Hardware Auto-negotiation Programming Sequence.....	19
2.6.8	Port ASIC Software Controlled Negotiation Programming Sequence.....	20
2.7	USXGMII Packet Control Header Implementation	21
2.8	PHY Implementation	26
2.8.1	Rx PHY Block	26
2.8.2	Tx PHY Block.....	27
2.8.3	MDI: Tx IPG and Pre-amble Processing on MDI.....	27
2.9	Electrical Specification	29
3	Appendix	31

Table of Figures

Figure 1: USXGMII Port ASIC Functional Block Diagram.....	7
Figure 2 USXGMII PHY Functional Block Diagram.....	8
Figure 3: Packet Control Header (PCH) Format	13
Figure 4: Packet Information Message Mapping in Pre-amble.....	21
Figure 5: Bit order for Serial CRC Computation	23
Figure 6: CRC Calculation	24

Table of Tables

Table 1: USXGMII Single Port Options	6
Table 2: Definition of channel control information passes between links.....	10
Table 3: Auto-neg Message aligned to Ordered Set followed Data and then by 7-bit Control Characters	11
Table 4 Auto-neg Message aligned to Control Code followed by Ordered Set.....	12
Table 5: Auto-neg Message aligned to Ordered Set and NO Control Codes.....	12
Table 6: Definition of channel control information passed between links via UsxgmiiPCH[47:0]	14
Table 7: UsxgmiiPCH Placement when SOP in LANE 0.....	25
Table 8: UsxgmiiPCH Placement when SOP in LANE 4 – Part1	25
Table 9: UsxgmiiPCH Placement when SOP in LANE 4 – Part 2	25

Definitions

MII - Media Independent Interface: A digital interface that provides a 4-bit wide datapath between a 10/100 Mbit/s PHY and a MAC sublayer. Since MII is a subset of GMII, in this document, we will use the term “GMII” to cover all of the specification regarding the MII interface.

GMII- Gigabit Media Independent Interface: A digital interface that provides an 8-bit wide datapath between a 1000 Mbit/s PHY and a MAC sublayer. It also supports the 4-bit wide MII interface as defined in the IEEE 802.3z specification. In this document, the term “GMII” covers all 10/100/1000 Mbit/s interface operations.

LPI- Low Power Idle: An alternative form of idle signaling that is used by the MAC to indicate that the PHY may enter a low power state and signal this change of state to the link partner; and is used by the PHY to signal to the MAC that the link partner has entered a low power state. The functions are defined by IEEE 802.3az in IEEE 802.3 clauses 22, 24, 25 (for 100Mb/s); 35, 36, 40, 70 (for 1Gb/s); 46, 48, 49, 55, 71, 72 (for 10GE); and 78 (for overall descriptions).

SGMII- Serial Gigabit Media Independent Interface: A digital interface that provides a 1.25 Gbps serial dual-data-rate datapath between a 1000 Mbit/s PHY and a MAC sublayer. Refer to ENG-46158

QSGMII- Quad Serial Gigabit Media Independent Interface: A digital interface that provides a 5.0 Gbps serial datapath between four 1000 Mbit/s PHY ports and a MAC sublayer. Refer to EDCS-540103

USGMII - Universal Serial Gigabit Media Independent Interface: A digital interface that provides capability to carry multi-port/multi-rate serial datapath between PHY ports and a MAC sublayer using 8B/10B coding. Refer to EDCS 1155168

USXGMII - Universal Serial 10 Gigabit Media Independent Interface: A digital interface that provides capability to carry multiport/multi-rate serial datapath between PHY ports and a MAC sublayer using 64B/66B coding.

1 Overview

USXGMII uses two data signals in each direction to convey frame data and link rate information between a single or multi-port PHY and the Ethernet MAC(s). This document specifies requirements for carrying multiple networks ports over a single PHY-MAC Interface. The maximum MAC/PHY SERDES speed is configured based on the maximum network port speed. Table 1, shows two options to support 100M/1G/2.5G/5G and 100M/1G/2.5G/5G/10G. A PHY/MAC can support appropriate USXGMII implementation to optimize system power and cost.

MAC-PHY IF Type	Number of Ports	Network Port Types	Replications – Lowest to Highest data speed	Port Mux/ Alignment Marker	PCS	SERDE S Speed (Gbps)
10G- USXGMII	1	100M/1G/2.5G/5G /10G	100, 10, 4, 2, 1	No	Clause 49	10.3125
5G- USXGMII	1	100M/1G/2.5/5G	50, 5, 2, 1	No	Clause 49	5.15625

Table 1: USXGMII Single Port Options

Other features:

- System Interface operates in full duplex mode only
- Ability to send PTP time stamp from PHY to MAC to improve accuracy/jitter on encrypted/non-encrypted PTP packet with MACSec is in the ASIC
- Hardware assisted auto-negotiation for all supported speeds
- Flexibility to add new features using Extension Field in the pre-amble

NOTES:

- USXGMII does not support 2.5G (2.578125Gbps) interface speed due to 2.5 replication requirement to carrying 1G. 2.5G is best suited for single port PHYs and recommendation is to use 2500BASE-X
- 10Mbps can be supported by adding appropriate replication

This document uses 10.3125Gbps SERDES to describe single network ports over single PHY-MAC interface, but same features applies to 5G unless stated otherwise.

Due to the high speed of operation, each of these signal pairs are realized as differential pairs thus optimizing signal integrity while minimizing system noise.

USXGMII leverages the 64B/66B PCS defined in IEEE 803.2ae Clause 49. The PCS is unchanged with additional functionality being added via the “ordered set” mechanism defined by IEEE.

Figure 1 and 2, illustrates high level view of Port ASIC and PHY with USXGMII interface respectively.

ASIC USXGMII block for 10G-USXGMII consists of

- Singel 10G and 100M/1G MACs. For 100M/1G GMII is mapped into XGMII in the Rate Adaptation/Replication block.
- There is a single PCS Clause 49 blocks, with additional ordered sets
- Auto-neg messages usign 16-bit configuration word
- 5.15625/10.325Gbps SERDES

NOTE: If the incoming data (Rx PHY) has less then 8-byte of pre-amble include SFD, the PHY MUST add appriorate bytes to ensure 7-bytes +SFD to meet XGMII/Clause 49 requirements. This applies to all data rates and PCH and non-PCH mode. See section 2.9 for more details.




Figure 1: USXGMII Port ASIC Functional Block Diagram

PHY USXGMII block for 10G-USXGMII consists of

- Single 10G and 100M/1G MACs. For 100M/1G GMII is mapped into XGMII in the Rate Adaptation/Replication block.
- There is a PCS Clause 49 blocks with additional ordered sets
- Auto-neg messages usign 16-bit configuration word
- 5.15625/10.325Gbps SERDES
- PHY PCS/PMA/PMD as appriorate for network interface type




Figure 2 USXGMII PHY Functional Block Diagram

When network is 100M/1GE mode, the traditional GMII data signals (TXD/RXD), data valid signals (TX_EN/RX_DV), and error signals (TX_ER/RX_ER) are encoded into XGMII words at a rate of 15.6254M/sec. For 2.5G, 5G and 10G mode XGMII is used.

1.1 In-band Control and Status Signaling

In-band signaling is carried via ordered sets, which are defined in IEEE 802.3ae Clause 46 (XGMII). Ordered sets provide the capability to carry messages of 3 bytes; two such messages are defined in 802.3ae -- *LocalFault* (0x00, 0x00, and 0x01) and *RemoteFault* (0x00, 0x00, and 0x02). These messages are used in USXGMII to signal faults within the channel carrying the message. The current implementation defines new ordered sets to support following features:

- Auto-negotiation Message - Ability to send auto-neg message from PHY to/from port ASIC- similar to SGMII auto-neg message.
- Packet Information Message: Ability to send Packet Type, SubportID, and Extension Field Type/Extension Field

1.1.1 Auto-neg Mechanism

MDI based auto-neg is performed based on the media type. For example for 10GBASE-T is based on IEEE 802.3 Clause 55 and Clause 28. Rx PHY block sends MDI auto-negotiated parameter to Auto-neg block.

USXGMII Auto-neg mechanism is based on Clause 37 (Figure 37-6) plus additional management control to select USXGMII mode. The PHY must provide a USXGMII enable control configuration via MDIO.

On power reset, USXGMII enable bit is de-asserted (logic “0”) and system interface on Port ASIC and PHY must assume normal XGMII/XFI (Clause 46/49) operation for 10Gbps. Once USXGMII enable bit is enabled via MDIO, auto-neg operation should follow Clause 37-6 functions with following modifications:

- an_sync_status=fail changed to block_lock=false (restart Autoneg FSM)
- Autoneg FSM will restart whenever the link changes
- rudi(invalid) changed to idle received during an_restart, ability_detect, acknowledge_detect
- Link_timer changed to be configurable from 1msec to 2 msec in steps of 0.1msec.
- Ability_match and acknowledge_match as per figure 37-6

1.1.2 Auto-negotiation Message

USXGMII also utilizes ordered sets to convey channelization and auto-negotiation information. The *UsxgmiiChannelInfo* message carries the status pertinent to the channel as per Table 2.

Bit	UsxgmiiChannelInfo[15:0] sent from the PHY to the MAC	UsxgmiiChannelInfo[15:0] sent from the MAC to the PHY	Default
15	Link: 1 = link up, 0 = link down	Same as PHY to MAC	0
14	Reserved for Auto-Negotiation acknowledge	1	0
13	0: Reserved for future use	0: Reserved for future use	0
12	Duplex mode: 1 = full duplex, 0 = half duplex	Duplex mode: 1 = full duplex, 0 = half duplex	1
11:9	Speed: Bit 11, 10, 9 : This should be simple network port speed 000 = 10Mbps 001 = 100 Mbps 010 = 1000 Mbps 011 = 10 Gbps 100 = 2.5 Gbps 101 = 5 Gbps 110 = Reserved 111 = Reserved	Speed: Bit 11, 10, 9: 000 = 10Mbps 001 = 100 Mbps 010 = 1000 Mbps 011 = 10 Gbps 100 = 2.5 Gbps 101 = 5 Gbps 110 = Reserved 111 = Reserved	010
8	EEE capability:1= supported, 0 = not supported	EEE capability:1= supported, 0 = not supported	0
7	EEE clock stop capability: 1= supported, 0 = not supported	EEE clock stop capability: 1= supported, 0 = not supported	0
6:1	0: Reserved for future use	0: Reserved for future use	0
0	Set to 1 (0 in SGMII)	1	1

Table 2: Definition of channel control information passes between links

Like SGMII, *UsxgmiiChannelInfo* uses a 1.6ms link timer. Any change in the status of the link requires the PHY to re-signal *UsxgmiiChannelInfo* words until message is received with Auto-negotiation acknowledge bit set to 1 or for the duration of the link timer.

At the XGMII interface, Auto-neg message uses the format below:

TXC [3:0]	Lane 0 [31:24]	Lane 1 [23:16]	Lane 2 [15:8]	Lane 3 [7:0]
0x1	Character Control Code=0x9C	Config[15:8]	Config[7:0]	Opcode for auto-neg = 0x03 (Cisco specific)

The above definition is compatible with the Clause 46 definition of sequence ordered sets. Cisco specific opcode is defined with a value of 0x03 in order to meet Lane 3 >= 0x03 specified in Table 46-5 of 802.3 Clause 46.

The tables below shows the possible insertion of the auto-neg message based on different position of ordered sets. The value of ordered set equal to 0x0, D3/D7 and 0x03 identifies the message as auto-neg message.

Input Data		Syn c	Payload (63:0)									
	Bit Position	0 1	65									
Data Block Format			Block Type									
O0 D1 D2 D3/C4 C5 C6 C7	10	0x4B	Config [15:8]	Config [7:0]	0x03	O4=0x0 (4-bit)	C4 (7-bit)	C5 (7-bits)	C6 (7-bit)	C7 (7-bit)		

Table 3: Auto-neg Message aligned to Ordered Set followed Data and then by 7-bit Control Characters

Input Data		Syn c	Payload (63:0)								
	Bit Position	0 1	2	65							
Data Block Format			Block Type								
Co C1 C2 C3/O4 D5 D6 D7		10	0x2D	C0 (7-bit)	C1 (7-bits)	C2 (7-bit)	C3 (7-bit)	O4=0x0 (4-bit)	Config [15:8]	Config [7:0]	0x03

Table 4 Auto-neg Message aligned to Control Code followed by Ordered Set

Input Data		Syn c	Payload (63:0)								
	Bit Position	0 1	2	65							
Data Block Format			Block Type								
O0 D1 D2 D3/O4 D5 D6 D7		10	0x55	Config[15:8]	Config[7:0]	0x03	O0 (4-bit)	O4=0x0 (4-bit)	Config [15:8]	Config [7:0]	0x03

Table 5: Auto-neg Message aligned to Ordered Set and NO Control Codes

1.1.3 Packet Control Header

The *UsxgmiiPCH* message conveys information that may be needed on a packet-by-packet basis. This message is optional, and is only needed when the Usxgmii instance is using features that require this message for PTP Time Stamps, MACsec etc.

PHY communicates with a port ASIC through Packet Control Header (PCH). PCH is 8 bytes and it replaces the preamble of the frame.

The details of PCH and fields are shown in *Figure 3*.




Figure 3: Packet Control Header (PCH) Format

Bit	UsgmiiPCH[47:0] sent from the PHY to the MAC	UsgmiiPCH[47:0] sent from the MAC to the PHY	Default
47:46	Packet type: 00: Ethernet Packet with PCH (packet information) 01: Ethernet packet, without PCH (packet information) 10: Idle Packet – Contains status data for a port – no packet data 11: Reserved	Packet type: 00: Ethernet Packet with PCH 01: Ethernet packet, without PCH (packet information) 10: Idle Packet – Contains status data for a port – no packet data 11: Reserved	0
45:42	Subport ID (Channel number) 3:0: Port 0 to 15	Subport ID (Channel number): MAC to PHY 3:0: Port 0 to 15	0
41:40	Extension Field Type 00: Ignore Extension Field 01: Extension Field contains 8-bit Reserved + 32-bit Timestamp 10: Reserved 11: Reserved	Extension Field Type 00: Ignore Extension Field 01: Extension Field contains Tag/Signature associated with for time-stamp to be taken by the PHY 10: Reserved 11: Reserved	0

Bit	UsgmiiPCH[47:0] sent from the PHY to the MAC	UsgmiiPCH[47:0] sent from the MAC to the PHY	Default
39:0	<p>Extension Field</p> <p>Extension Type = 00: Extension Field is ignored</p> <p>Extension Type = 01:</p> <ul style="list-style-type: none"> • 39:32: Reserved. • 31:0: PTP Timestamp <p>Extension Type = 10 and 11 are reserved</p>	<p>Extension Field</p> <p>Extension Type = 00: Extension Field is ignored</p> <p>Extension Type = 01:</p> <ul style="list-style-type: none"> • 39:32: Reserved • 31:16: Reserved • 15:0 Signature associated with the egress PTP Timestamp to be done in PHY <p>Extension Type = 10 and 11 are reserved:</p>	Refer to PHY spec for full details

Table 6: Definition of channel control information passed between links via UsxmiiPCH[47:0]

Six byte UsxmiiPCH message is sent in the pre-amble of the packet. It is inserted between SOP (0xfb) and Header CRC. Refer to section 2.6 for more details.

2 Implementation Specification

This section describes summary of MII, GMII, XGMII byte replication, mapping to 64b data + 8b control USXGMII words for translation into a USXGMII stream via the Clause 49 PCS.,.

2.1 XGMII Mapping

XGMII naturally maps very cleanly into USXGMII words. Packet SOPs are already aligned into lane 0 or 4.

UsxgmiiChannelInfo is required to convey speed negotiation between PHY and MAC.

UsxgmiiPCH is required for feature using the Extension Field for features such as passing PTP Time Stamp between PHY and MAC Device. Extension Field could also be used for additional feature in the future.

2.2 GMII Mapping

GMII, being a byte-wide interface, requires additional translation to map into USXGMII words. The SOP Adjust feature is required. During USXGMII encoding, when an SOP is generated, if it is not naturally aligned on lane 0 or 4, it is left-shifted (i.e. made “earlier”) to the nearest legal SOP lane.

NOTE: If the incoming data (Rx PHY) has less than 8-byte of pre-amble include SFD (GMII), the PHY MUST add appropriate bytes to ensure 7-bytes +SFD to meet XGMII/Clause 49 requirements. This applies to all data rates and PCH and non-PCH mode. Refer to section 2.8 for more details.

2.3 MII Mapping

Same as GMII

2.4 Pause Frame Support

There is no PAUSE frame support at USXGMII level, port level pause frames are send over USXGMII interface as payload. Channel level pause frame behavior should be as per 802.3 Annex 31B. Flow control is done using *UsxgmiiPacketInfo* as described above.

2.5 Auto-neg Mechanism

MDI based auto-neg is performed based on the media type. For example for 10GBASE-T is based on IEEE 802.3 Clause 55 and Clause 28. Rx PHY block sends MDI auto-negotiated parameter to Auto-neg block.

USXGMII Auto-neg mechanism is based on Clause 37 (Figure 37-6) plus additional management control to select USXGMII mode. The PHY must provide a USXGMII enable control configuration via MDIO.

On power reset, USXGMII enable bit is de-asserted (logic “0”) and system interface on Port ASIC and PHY must assume normal XGMII/XFI (Clause 46/49) operation for 10Gbps. Once USXGMII enable bit is enabled via MDIO, auto-neg operation should follow Clause 37-6 functions with following modifications:

- an_sync_status=fail changed to block_lock=false (restart Autoneg FSM)
- Autoneg FSM will restart whenever the link changes
- rudi.invalid changed to idle received during an_restart, ability_detect, acknowledge_detect
- Link_timer changed to be configurable from 1msec to 2 msec in steps of 0.1msec.
- Ability_match and acknowledge_match as per figure 37-6

2.5.1 Transmitting Configuration Words

The transmit configuration words are enabled by software. Configuration words are transmitted during hardware auto-negotiation or software controlled negotiation on USXGMII interface.

During hardware auto-negotiation (see section 1.1.2) the configuration words are transmitted in accordance to IEEE 802.3 Clause 37 auto-negotiation. During auto-negotiation, the packet data or idles on the MII interface of Tx MAC are discarded and not sent to Tx PCS. Instead, the auto-negotiation block drives the configuration words to Tx PCS on the MII interface. However if Tx MAC transmits errors or link faults, the auto-negotiation is interrupted and error/link fault information is passed on the MII interface from Tx MAC to Tx PCS. Other than data/idles all other control words are pass-through e.g. errors and link faults.

Once auto-negotiation is completed, the Tx MAC output drives the Tx PCS using the MII interface. USXGMII Auto-neg is NOT re-started due to error/link fault conditions, Auto-neg is started due to network interface link changes or forced via software control.

Software may choose to use software controlled negotiation (section 2.5.7) instead of hardware auto-negotiation. This allows for diagnostic debug or “force link speed” under software control.

2.6 Rate Adaptation - Replicating Transmit Bytes

Once software programs the link speed of 10M, 100M, 1G, 2.5G and 5G the logic replicates the 4-byte words on the transmit MII interface based on 10G and 5G SERDES modes in the following manner.

2.6.1 Tx 10G Mode

- SOP word (4-bytes) is transmitted only once; in remainder of the 1 (5G), 3(2.5G) 9 (1G) or 99 (100M) 4-byte words are replaced by 0xAA. This is compatible with the checks of the Tx PCS state machine that back to back SOPs are not allowed.
- Each word of the preamble and payload is replicated 2/4/10/100 times for 5G/2.5G/1G/100M respectively
- EOP is transmitted once and the remaining 1/3/9/99 bytes are idles. This is compatible with the checks of the Tx PCS state machine that back to back EOPs are not allowed.
- Each idle byte is replicated 2/4/10/100 times. This ensures that the data rate in case of 2/4/10/100 does not exceed the frame rate.
- All other sequence ordered sets and control sets are replicated 2/4/10/100 times.

Below is an example showing replication of a transmit frame on the 4-byte MII interface.

Transmit frame:

```
fb555555 555555d5 01020304 05060708 ... fd070707
```

Tx MII Interface Data [31:0] view of the transmit frame:

```
55555fb d5555555 04030201 08070605 ... 070707fd
```

Replicated pkt on the 4-byte MII Data bus [31:0] for 1G mode (x10 replication):

SOP: 55555fb 555555aa 555555aa 555555aa 555555aa 555555aa 555555aa 555555aa
555555aa 555555aa

Preamble: d5555555 d5555555 d5555555 d5555555 d5555555 d5555555 d5555555
d5555555 d5555555 d5555555

Data: 04030201 04030201 04030201 04030201 04030201 04030201 04030201 04030201
04030201 04030201
08070605 08070605 08070605 08070605 08070605 08070605 08070605 08070605
08070605 08070605
.....

EOP: 070707fd 07070707 07070707 07070707 07070707 07070707 07070707 07070707
07070707 07070707

2.6.2 10G Mode Sampling Received Bytes

The logic receives SOP and depending on the software-programmed or hardware auto-negotiated link speed, samples every 2th, 4th, 10th and 100th subsequent 4-byte word to form the Rx MII data for 5G, 2.5G, 1G and 100M network port speed respectively. This is then passed to Rx MAC which processes the packet, checks CRC and other errors and ultimately passes the frame to user interface.

Any configuration words, link fault words or errors received are similarly sampled at the MII interface and the 2th, 4th, 10th and 100th 4-byte word is passed to the Rx MAC for 5G, 2.5G, 1G and 100M network port speed respectively.

2.6.3 Tx 5G Mode

- SOP word (4-bytes) is transmitted only once; in remainder of the 1 (2.5G), 4(1G) or 49 (100M) 4-byte words are replaced by 0xAA. This is compatible with the checks of the Tx PCS state machine that back to back SOPs are not allowed.
- Each word of the preamble and payload is replicated 2/5/50/500 times for 2.5G/1G/100M respectively
- EOP is transmitted once and the remaining 1/3/9/99/999 bytes are idles. This is compatible with the checks of the Tx PCS state machine that back to back EOPs are not allowed.
- Each idle byte is replicated 2/4/10/100/1000 times. This ensures that the data rate in case of 2/4/10/100/1000 does not exceed the frame rate.
- All other sequence ordered sets and control sets are replicated 2/4/10/100/1000 times.

Below is an example showing replication of a transmit frame on the 4-byte MII interface.

Transmit frame:

```
fb555555 555555d5 01020304 05060708 ... fd070707
```

Tx MII Interface Data [31:0] view of the transmit frame:

```
555555fb d5555555 04030201 08070605 ... 070707fd
```

Replicated pkt on the 4-byte MII Data bus [31:0] for 1G mode (x5 replication):

SOP: 555555fb 555555aa 555555aa 555555aa 555555aa

Preamble: d5555555 d5555555 d5555555 d5555555 d5555555

Data: 04030201 04030201 04030201 04030201 04030201
08070605 08070605 08070605 08070605 08070605

.....

EOP: 070707fd 07070707 07070707 07070707 07070707

2.6.4 Rx 5G Sampling Received Bytes

The logic receives SOP and depending on the software-programmed or hardware auto-negotiated link speed, samples every 2th, 5th and 50th subsequent 4-byte word to form the Rx MII data for 2.5G, 1G and 100M network port speed respectively. This is then passed to Rx MAC which processes the packet, checks CRC and other errors and ultimately passes the frame to user interface.

Any configuration words, link fault words or errors received are similarly sampled at the MII interface and the 2th, 5th and 50th 4-byte word is passed to the Rx MAC for 2.5G, 1G and 100M network port speed respectively

For other replication in other USXGMII-M speed, refer to Table 1.

2.6.5 Receiving Configuration Words

When valid configuration words are received, the logic check that the last 3 consecutive configuration words received are identical and signals an interrupt for software to read the received configuration word.

2.6.6 Clocking

The auto-negotiation block runs on PCS/MAC clock.

2.6.7 Hardware Auto-negotiation Programming Sequence

Software will follow the programming sequence below to achieve auto-negotiation with the link partner. If hardware auto-negotiation doesn't converge to a link speed (step 8 below), software may timeout and restart hardware auto-negotiation. Following uses 10G as an example.

1. Reset.
2. 10G link up. At this time the link is active at 10G, no replication or sampling is being done. Device can receive and transmit packets on the native 10G link.
3. Programs the link timer registers, if desired. Otherwise the default value is used.
4. Programs transmit configuration word.
5. Enables hardware auto-negotiation and disables software negotiation in a register.
6. Device starts transmitting configuration words in accordance to the auto-negotiation state diagram of IEEE 802.3z Clause 37.
7. Device transitions through the auto-negotiation states in hardware, without any software involvement. Since the link partner may start transmitting configuration words before, the reception and transmission of configuration words happen independently.

8. Upon completion of auto-negotiation, hardware sets the link speed register. At this point Device starts transmitting the replicated idles and packets and sampling the received data. Upon receiving SOP, Device recalibrates the sampling start point.
9. If 10G link is lost or regained, the software is expected to disable auto-negotiation and re-enable auto-negotiation.

2.6.8 Port ASIC Software Controlled Negotiation Programming Sequence

If software wishes to disable hardware auto-negotiation and manage the negotiation process itself, it follows the sequence below.

If software wishes to read the PHY directly (using MDIO, for example) it may skip steps 4 through 6 below and program the Link Status register directly.

1. Reset.
2. 10G link up. At this time the link is active at 10G/5G, no replication or sampling is being done. Device can receive and transmit packets on the native 10G/5G link.
3. Disable hardware auto-negotiation and enables software negotiation in Auto-neg Control register.
4. Programs transmit configuration word in Tx Config Control register and enables transmission of configuration words.
5. Check for the valid bit of Rx Config Word register. When set, it reads the configuration word received.
6. Enable transmission of idles in Tx Config Control register.
7. Programs the link speed and negotiation complete fields in Link Speed Control Status register. At this point Port ASIC starts transmitting the replicated idles and packets and sampling the received data. Upon receiving SOP, Port ASIC recalibrates the sampling start point.
8. If 10G/5G link is lost or regained, the software is expected to restart the negotiation sequence.

2.7 USXGMII Packet Control Header Implementation

When Packet Control Header is used, the PHY and ASIC must use fixed offset of 7-bytes from 0xfb to match 0xd5 on the 8th byte to detect start of packet data. The data after 0xfb and before 0xd5 is the Packet Information Message and packet data is after 0xd5.




Figure 4: Packet Information Message Mapping in Pre-amble

2.7.1.1 Start of Packet (SOP)

The SOP symbol indicates the start of packet. The value of SOP MUST be 0xFB (Hex) as per IEEE 802.3 standard.

2.7.1.2 Packet Type

Packet Type identifies type of packet:

- 00: Ethernet Packet with status header
- 01: Ethernet packet, no Status Header (packet information)
- 10: Idle Packet – Contains status data for a port – no packet data
- 11: Unused (Reserved)

2.7.1.3 Sub-portID

Sub-portID is provided as additional information to be used at layer protocol layers. For single port is set to 0.

2.7.1.4 Extension Field Type

This 2-bit field type defines the content of Extension Field. Refer to appropriate PHY for the definition.

2.7.1.5 Extension Field

This 40-bit Extension Filed carries data defined by Extension Field Type, e.g. Time-Stamp.


2.7.1.6 CRC computation

Header CRC is computed using CRC-8 algorithm with polynomial x^8+x^2+x+1 over the 6-bytes header data (inclusive of Packet Type to Extension Field). The initial value of CRC-8 computation is 0. The CRC remainder is exclusive-or'ed with 0x55 to header.crc value. The result is copied to the header.crc field on header generation and compared to header.crc on header reception. Header CRC generation is based on ITU-T I.432

Readers must pay attention to the differences in bit numbering conventions and bit [Page 15]. CDL Working Group CDL Specification 17 May 2002 transmission order between ITU-T standard I.432 and this specification.

The I.432 standard numbers the LSB of a multi-bit field as zero and bits within a byte are transmitted MSB first. See the Section "Definitions" for bit numbering conventions and bit transmission order applicable to this specification.

The Header CRC is an 8-bit sequence calculated over 6 bytes PCH but excluding the SOP symbol and the Header CRC. The 48-bits long relevant portion of the CDL header is taken to represent a polynomial of order 47. The coefficients can only have the value 0 or 1. The least significant bit of the first byte of the header represents the coefficient of the highest order (x^{47}) term. The polynomial operations are performed modulo-2.

**Figure 5: Bit order for Serial CRC Computation**

Other examples for CRC:

PCH[47:0] = 0x2910_4602_7710. PCHCRC = 0x0B
 PCH[47:0] = 0x2910_4602_7720. PCHCRC = 0x07
 PCH[47:0] = 0x2910_4602_7730. PCHCRC = 0x0F
 PCH[47:0] = 0x2910_4602_7740. PCHCRC = 0x01

The Header CRC MUST be recomputed whenever any of the fields in the header is changed and MUST be passed transparently whenever the fields of the header do not change.
 The receiver in to-CDL-network direction must perform HCRC and maintain count of packets with failed HCRC for performance monitoring purposes. CDL information MUST NOT be extracted from, made use of or inserted into packets that failed HCRC. Such packets MAY also be discarded by the transmitter in from-CDL-network direction.

The linear feedback shift register below can be used to calculate the 8-bit CRC in a bit serial fashion. 8, 16 or 32 bit parallel implementations can be deduced from the bit serial implementation.

NOTE: The feedback on this implementation is different than 32-bit frame CRC and Clause 55




Figure 6: CRC Calculation

2.7.1.7 Mapping Packet Information Message to 64B/66B PCS Layer

UsxgmiiPCH message is carried in the preamble of the SOP.

Input Data		Sync	Payload (63:0)										
	Bit Position	0 1	2	65									
Data Block Format			Block Type										
S0 D1 D2 D3/D4 D5 D6 D7	10	0x78	<i>UsxgmiiPCH(47:0)</i>										CRC (8bit)

Table 7: UsxgmiiPCH Placement when SOP in LANE 0

Input Data		Sync	Payload (63:0)										
	Bit Position	0 1	2	65									
Data Block Format			Block Type										
O0 D1 D2 D3/S4 D5 D6 D7	10	0x66	0x00	0x00	0x00	O0	O4	<i>UsxgmiiPCH(47:24)</i>					

Table 8: UsxgmiiPCH Placement when SOP in LANE 4 – Part1

Input Data		Sync	Payload (63:0)										
	Bit Position	0 1	2	65									
Data Block Format			Block Type										
D0D1D2D3D4D5 D6D7	0 1	<i>UsxgmiiPCH (23:0)</i>			CRC (8-bit)	Pkt Data0	Pkt Data1	Pkt Data2	Pkt Data3				

Table 9: UsxgmiiPCH Placement when SOP in LANE 4 – Part 2

2.8 PHY Implementation

This is vendor dependent and may choose any implementation as long as the features are supported described above. Figure 2 show major Rx and Tx blocks required in PHY to support USXGMII features.

2.8.1 Rx PHY Block

Rx PHY consists of:

- Media Dependent Interface – 802.3an for Copper
- Tx Rate Adaptation – Rate adaptation for 100M/1G/2.5G and or 5G, UsxgmiiPCH and GMII to XGMII mapping
- Tx PCS – 802.3 – 2008 Clause 49 (10GBASE-R) with modifications
- Tx System Interface - PMA/PMD SERDES (10.3125Gbps)

For 10G/5G/2.5G data rates, MDI interface provides XGMII (Clause 46) interface to Tx Rate Adaptation block, while in 100/1000M rate, data/control information is provided over GMII (Clause 35).

Refer to section 2.5.2 for replications/deletions for each MAC-PHY rate based on network port combinations.

NOTES:

- *When MDI is in auto-neg, PHY to send data from user programmable register. Data could be IDLE or Local Fault.*
- *When txer=1 and txen=1 from mii copper receiver interface, xgmii will send control character /E/ which is 0xfe*
- *Carrier extension for 1GE is not supported (No Half-Duplex support)*

2.8.1.1 MDI: Receive IPG and Pre-amble Processing on MDI

PHY must be able to meet IPG and Pre-ambles received on the MDI based on the speed:

- 10M/100M/1G/2.5G/5G with Clause 36 (8b/10b coding) – USXGMII PHY is required to receive Ethernet packets from remote PHY using 8b/10b PCS coding with IPG of 8-bytes and 7-byte (6 + SFD) Preamble as per IEEE 802.3. But legacy 10M/100M/1G implementation are not always following IEEE 802.3 IPG and pre-amble rules, and hence

USXGMII PHY needs to be more flexible in terms of receiver design. It is possible the legacy equipment may send few packets with IPG less than 8-bytes and Preamble of more than 8-bytes. USXGMII PHY should avoid exact match of 8-byte IPG and 8-byte of pre-amble. This can cause overflow of the PHY which should be handled gracefully and provide overflow status accessible via MDIO/management interface

- 802.3Clause 46/49 and 126: PHY must be able to receive minimum of 5-byte IPG (4-byte Idle + T) and pre-amble of 8-bytes

2.8.1.2 USXGMII Transmit to MAC (ASIC)

PHY must be able to receive a packet with minimum of 5-byte IPG. For debug, it's recommended that the PHY detect number of IPG bytes received and maintain overflow status of its transmit buffer.

2.8.2 Tx PHY Block

Tx PHY consists of (from Port ASIC to PHY):

- Rx System Interface - PMA/PMD SERDES (10.3125Gbps)
- Rx PCS – Clause 49 (10GBASE-R) with modifications
- Rx Rate Adaptation – Rate adaptation for 10/100/1000M and XGMII to GMII mapping
- Media Dependent Interface – e.g. IEEE 802.3 clause 55 for Copper

Phy's Rx System Interface from Port ASIC receives data using clause 49 as shown in figure 1. Data from Rx PCS is USXGMII format. In 100M/1000M, Rx Rate Adaptation removes replicated data and regenerates GMII data/control information with preamble, SFD and Carrier Extension as appropriate. In 10G mode, XGMII is sent to PHY.

When in auto-neg or fault status, Rx Adaptation block removes special sequence ordered-set and forward them to USXGMII Auto-neg block. Only the valid data/control Clause 35/46 is sent to PHY MDI block. Refer to section 2.5.2 for replications/deletions for each MAC-PHY rate based on network port combinations.

Rx Rate Adaptation passes the Rx –configuration/error status to Auto-neg block for clause 37 Rx processing.

2.8.2.1 USXGMII Receive from MAC (ASIC)

PHY must be able to receive a packet with minimum of 5-byte IPG from the MAC. For debug, it's recommended the PHY to detect number of IPG bytes received and maintain overflow status.

2.8.3 MDI: Tx IPG and Pre-amble Processing on MDI

PHY must be able to meet IPG and Pre-ambles received on the MDI based on the speed:

- 802.3bz Clause 46/49 and 126: PHY must send minimum IPG of 8-bytes with average of 12-bytes. Transmit Pre-amble must be 8-bytes

2.9 Electrical Specification

CML, (Current Mode Logic) is by far the most common Serdes IO standard in use today. The signal swing provided by the CML output is small, resulting in low power consumption. The driver and receiver are often self-terminated, eliminating external components and minimizing transmission line impedance discontinuity effects on timing and signal integrity.

The following section details the requirements for the high speed electrical interface that will operate at 1.25Gsym/s or 10Gsym/s using NRZ coding (hence 1 bit per symbol at electrical level). Connections are point to point balanced differential pair with 100 Ohm nominal differential impedance and signaling is unidirectional. Clock and data are embedded hence CDR is required in the receiver. The link should operate with a BER of 10^{-15} . It supports both AC and DC coupled operation. However, DC coupling of PHY to MAC is required since it optimizes system cost, complexity, and signal integrity.

The following section details the requirements for the USXGMII – 10 Gbps electrical interface

- IP vendor must provide AMI models and support internal eye (ability to measure the eye at the RX sampling latch)
- IP vendor must provide Vertical and Horizontal Eye opening at a BER (used as a PASS/FAIL criteria for AMI simulation and internal eye diagnostic)
- 10Gsym/s using NRZ coding (hence 1 bit per symbol at electrical level)
- Connections are point to point balanced differential pair
- 100 Ohm nominal differential impedance
- Unidirectional signaling on each differential pair
- Clock and data are embedded hence CDR is required in the receiver
- The link should operate with a BER of 10^{-15}
- AC coupling required between driver and receiver

The USXGMII 10 Gbps electrical specification will be the same as the 10GBASE-KR electrical characteristics as defined in section 72.7 and Annex 69B of the IEEE 802.3-2008

Few exceptions will apply: This is to allow low cost interface implementation, while meeting system requirements.

1. Maximum insertion loss is 15dB @ 5 GHz (10 Gbps) (vs ~25dB based on the KR spec). This is to allow low cost interface implementation, while meeting system requirements.




Figure 69B-5—Insertion loss limit for 10GBASE-KR

2. Maximum ICR at 5GHz is 15dB (vs ~25dB)




Figure 69B-8—Insertion loss to crosstalk ratio limit

$$ICR(f) = -IL(f) + PSXT(f)$$

3 Appendix

Following pseudo code is provided to clarify bit order in computation of CRC-8. This used be used with reference to section 2.6.16. This example is included for reference/validation only – actual implementation is left for the implementer.

```

logic [47:40] pchHdr;
logic [63:0] preamble;
logic [7:0] pchCrc;

// Preamble contains SOP, PCH, PCH-CRC
assign preamble = {8'hfb, pchHdr, pchCrc}; // SOP + PCH (6-bytes) + PCH-
// CRC (1-byte)

// Generate PCH-CRC
Pch48bCrc8 pch48bCrc8
(
    .sysClk      (sysClk),
    .sysClkReset (sysClkReset),
    .data        ({pchHdr[7:0], pchHdr[15:8], pchHdr[23:16],
    pchHdr[31:24], pchHdr[39:32], pchHdr[47:40]}),
    .crc8        (pchCrc)
);

// 
// PCH CRC Generator/Checker
//
module NifPch48bCrc8 (
    input logic          sysClk,
    input logic          sysClkReset,
    input logic [5:0][7:0] data,    // 48-bit PCH: data[0][7:0] is the
                                    // first PCH byte rcvd/transmitted
    output logic [7:0]    crc8
);

parameter      crcPolynomial = 8'b10000011; // per USXGMII spec:
x8+x2+x1+1
integer        pchByteNum, bitNum;

logic [5:0][7:0] feedbk;
logic [5:0][7:0] transTerm;
logic [5:0][7:0] nextCrc;
logic [5:0]       dataXorBitIn;

// 8-bit CRC output
assign crc8 = nextCrc[5] ^ 8'h55;    // per USXGMII spec

always_comb begin

```

```

        for (pchByteNum=0; pchByteNum < 6; pchByteNum=pchByteNum+1) begin
// for each of the 6 PCH bytes

        if (pchByteNum==0)
            nextCrc[pchByteNum] = 8'h00;
        else
            nextCrc[pchByteNum] = nextCrc[pchByteNum-1];

            for (bitNum=0; bitNum < 8; bitNum=bitNum+1) begin           //
for each bit in a PCH byte
            feedbk[pchByteNum]      = {8{nextCrc[pchByteNum][7]}};
// feedback term
            transTerm[pchByteNum]   = (feedbk[pchByteNum] &
crcPolynomial) ^ nextCrc[pchByteNum];
            dataXorBitIn[pchByteNum] = data[pchByteNum][bitNum] ^

nextCrc[pchByteNum][7]; // LSB in first
            nextCrc[pchByteNum]     = {transTerm[pchByteNum][6:0],
dataXorBitIn[pchByteNum]};
            end
        end
    end

endmodule

```

